

Ship faster. Without breaking production.

Progressive delivery, targeting, experimentation, and AI control — on one platform, with an instant undo at every step. You stop choosing between speed and safety.

Prepared for ABC Company



PART 1

Release & Remediate

Deploy ≠ release. And the undo button is instant.

VALUE

- Ship code dark; release as a business decision — no redeploy.
- Instant kill switch: roll back in seconds, not a rebuild.
- A trigger lets a monitoring alert remediate without engineers.

Proof: Ally — 97% fewer weekend releases · Jackpocket — 90% fewer deploy incidents · AlayaCare — 50% lower MTTR.

“Isn't LD now a prod single point of failure?” SDKs evaluate locally & cache last-known rules; Relay Proxy keeps it in-VPC; 99.99% SLA.



PART 2

Target

Test in production, safely — to exactly the right audience.

VALUE

- Progressive rollout: you → team → 1% → enterprise, in real prod.
- Individual + rule-based targeting, no code change.
- Multi-context: target the user AND their org / plan / device.

Proof: Savage X Fenty — 15% site-performance gain · Blue Cross of Idaho — 75% less portal troubleshooting.

“At 40k/day, do flags add latency? And our PII?” In-memory local evaluation — no per-request call. Anonymous contexts + private attributes never leave your app.



EXTRA CREDIT

Experimentation

The flag you release with is the flag you measure with.

VALUE

- No second SDK — the flag is the experiment.
- Tie variations to conversion, revenue, retention.
- PMs self-serve; ship the winner with the same flag.

Proof: SmugMug — PMs A/B test with no engineer time · Gamma — 20% lift in paid subscriptions.

“We already use Optimizely / Statsig / GA.” Those bolt on. Here the experiment runs on the flag you already ship — one platform, no duplicate instrumentation.



EXTRA CREDIT

AI Configs

Prompt and model become runtime config — not a redeploy.

VALUE

- Change model, prompt, or temperature live — no redeploy.
- Target & canary AI like any feature; experiment on prompts.
- Instant rollback of a bad prompt or runaway model.

Proof: Dior — AI change time from ~15 min to instant · Hireology — tests prompt/model variants in seconds.

“How is this different from LangSmith / a prompt tool?” Same platform as your flags, targeting & experiments — model + prompt governed and measured together, in prod.